

Technology fields, automation and task reinstatement

Working paper — not for circulation

Joakim Storck and Jonatan Andersson

Draft: June 11, 2026

Abstract

We use the geometric task-space framework of Storck and Andersson (2026) to study how the reach of a technology determines whether task creation benefits or harms workers in different parts of the labour market. The framework places occupations as task bundles in a two-dimensional task space calibrated from occupational embeddings. The market price of skill is a field over this space, directional rather than uniform, so that depth is rewarded in the socio-cognitive directions and not in the service and manual directions. Automation removes the high-priced tasks from a bundle; reinstatement returns a fraction γ of the displaced mass as new tasks, which attach to an occupation only where the priced capability they require is already held. Occupations are therefore redefined rather than abolished, and their wages move with the priced content of their bundles. New tasks that no occupation can attach remain unbound, and their concentration marks where new occupations may emerge.

We calibrate the technology parameters against task-level AI exposure scores from Gmyrek et al. (2025) and test the model’s central spatial prediction against Bureau of Labor Statistics occupational employment data for 741 occupations over 2019–2024. A difference-in-differences specification exploiting the ChatGPT launch of November 2022 shows that occupations at the gradient boundary of the augmentation field experienced significantly higher employment growth in the post-2022 period, while occupations in the field’s core experienced lower growth. The boundary zone effect is robust across six specifications and holds in an event study with clean pre-trends.

The model also reveals a structural threshold at $z_K \approx 0.38$ – 0.40 . Below this threshold, reinstatement concentrates new tasks in the cognitive and investigative sectors and harms manual workers through general equilibrium redistribution. Above it, the technology’s gradient boundary reaches the manual arc and reinstatement benefits manual workers in employment, though the directional price of skill mutes the wage benefit in those directions. This spatial condition, the reach of the technology relative to the distance between the technology centre and the manual arc, cannot be derived from a one-dimensional task-space model.

Contents

1	Introduction	3
2	The Spatial Occupational Geometry	4
3	Model	4
3.1	Overview	4
3.2	The Capability Plane and the Price of Skill	5
3.3	Tasks, Occupations, and Wages	7
3.4	Technology and the Operated Regime	7
3.5	Reinstatement: The Bundle Operator	9
3.6	Unbound Tasks and Candidate Occupations	12
3.7	Employment and Mobility	12
3.8	Augmentation and the Character of the Technology	14
4	Empirical Implementation and Calibration	15
4.1	Data and frozen inputs	15
4.2	The capability plane	16
4.3	The price field	17
4.4	From centroids to bundles	18
4.5	The price functional and the gate weights	18
4.6	The wage wedge	19
4.7	Technology calibration	19
5	Illustrative Cases	19
6	Empirical Analysis	19
7	Discussion	20

1 Introduction

The division of labor segments the labor market into specialized domains. Workers accumulate human capital specific to their occupation or sector, and technologies are adopted unevenly across the task distribution. When a technology displaces workers from one set of tasks, the consequences for wages and employment depend on whether new tasks are created to absorb the displaced labor: the reinstatement mechanism of Acemoglu and Restrepo (2019). In the canonical one-dimensional framework, reinstatement unambiguously raises the labor share because new tasks are by construction assigned to labor at the boundary of the technology’s operating domain (Acemoglu and Restrepo 2018). The question of which workers benefit is not well posed in a model where labor enters as two or three aggregate skill groups distinguished only by their position on a single task index.

This paper develops a spatial generalization of the task-based framework and tests its central prediction against occupational employment data. The framework places occupations as task bundles in a two-dimensional task space calibrated from semantic embeddings of O*NET task statements (Storck and Andersson 2026). The two dimensions capture domain orientation and specialization depth, placing occupations in a continuous space that reflects the actual content of work rather than predefined skill categories. The market price of skill is a field over this space, directional rather than uniform. Technologies are represented as augmentation fields: Gaussian productivity enhancements centered at a point in task space with a radius that determines their reach across occupational domains.

The spatial structure generates a precise prediction about where new tasks appear. The gradient of the augmentation field is maximized at distance z_K from the technology center: at the boundary between the core, where the technology dominates and the marginal value of complementary human activity is low, and the periphery, where the technology is irrelevant. New tasks that validate, coordinate, and integrate automated outputs emerge precisely in this boundary zone. Whether reinstatement benefits workers therefore depends not on the reinstatement elasticity alone, but on where the gradient boundary falls relative to the spatial ownership structure of task space, and on the price the market attaches to work in that direction. A broad, inexpensive technology centered in the cognitive arc can generate task creation that harms manual workers through general equilibrium redistribution rather than offsetting their displacement.

We test this prediction using Bureau of Labor Statistics Occupational Employment and Wage Statistics (OEWS) data for 741 occupations over 2019–2024, combined with task-level AI exposure scores from Gmyrek et al. (2025) mapped onto the occupational geometry. The treatment variable is not binary but continuous: for each occupation we compute both its exposure to the augmentation field and its position relative to the field’s gradient. A difference-in-differences specification exploits the ChatGPT launch of November 2022 as

an exogenous shift in the technology’s reach. Occupations at the gradient boundary show significantly higher employment growth in the post-2022 period relative to pre-trends, while occupations in the field’s core show lower growth. The boundary zone effect is robust across six specifications and holds in an event study with clean pre-trends.

The paper proceeds as follows. Section 2 introduces the spatial occupational geometry. Section 3 presents the task-based production framework and its spatial extension, including the price of skill, the augmentation field, displacement, and displacement-anchored reinstatement through a bundle operator. Section 4 describes the calibration of technology parameters against empirical AI exposure data. Section 5 illustrates the model with two scenarios drawn from the Acemoglu–Restrepo taxonomy. Section 6 presents the empirical analysis. Section 7 discusses implications and limitations.

2 The Spatial Occupational Geometry

[TODO: Kort presentation av geometrins från Paper 1: taskspace, vinkeldimension (domän), radiedimension (specialisering), occupational centroids och kompetensradier. En figur (taskspace-kartan). Ca 1–1.5 sidor.]

The value of this representation can be shown through its ability to illustrate successive waves of labor-replacing technology. The routine-task automation studied by Autor, Levy, and Murnane (2003) targeted occupations near the Technical-Physical pole, contributing to the secular decline of middle-wage manufacturing employment (Acemoglu and Restrepo 2020; Graetz and Michaels 2018; De Vries et al. 2020). Generative AI capabilities (Hampole et al. 2025; Eloundou et al. 2024) extend into parts of the Social-Cognitive region previously considered safe, raising concerns about the exposure of college-educated work. Each technology can be characterized by the region of the task space it affects, and each occupation by its position relative to that region. This is a question the geometry is built to address but that a partitioned or one-dimensional model cannot express.

3 Model

We embed the geometry into a production framework that builds on the canonical task-based model of Acemoglu and Restrepo (2019), extending it by replacing abstract factor categories with a continuous task space and by letting technology redraw the task bundles that define occupations.

3.1 Overview

The model has two layers. A *task layer* is a set of fields over the task disk. It carries production, the technology, displacement, reinstatement, and the creation of unbound

tasks. A *worker layer* carries employment, mobility, and labor supply. The two layers are linked by the bundles that define occupations, and it is these bundles that automation and reinstatement rewrite.

The task layer is family-free. Job families are an external classification by industrial context, and the geometry of Storck and Andersson (2026) is built without such categories. We work at the level of tasks, the level at which the geometry is constructed and at which technology exposure is measured. Occupations enter as bundles of tasks, not as primitives.

Wages are pinned by the price of skill and the task bundles alone, and the theory layer carries no occupation-specific wage parameter. The price of skill is a field over the disk, identified by Storck and Andersson (2026) and carried here as the price of labor at each task location. What the field leaves unexplained is recorded in a measurement layer: a family-level wage wedge η_g , measured once on the cross-section and held fixed under technology shocks. The wedge is not part of the theory; simulations are reported with and without it (Section 4).

The model’s objects accordingly divide into three layers with different epistemic status. The *theory* is two objects: the capability plane and the price functional (Section 3.2). The *measurement* is their frozen coordinates: the plane coefficients, the components of the price functional, and the wage wedge, each estimated once on the cross-section and never free under shocks. The *simulation* owns the free primitives: the technology (p_K, z_K, A_K, s_K) and the economy ($R, \tau, \gamma, \ell, \beta, c, \kappa$). A technology scenario varies the last list only.

3.2 The Capability Plane and the Price of Skill

Storck and Andersson (2026) shows that the labor market does not value position in task space isotropically. The return to radial depth is positive toward the analytical directions and absent toward the service and manual directions, and what the market prices is the socio-cognitive cluster of capabilities concentrated in the northern half.

These findings are not independent, and the model rests on the structure that unifies them. Storck and Andersson (2026) establishes four regularities about how capabilities organize over the disk: the mean similarity of capability deviation profiles between occupations follows $a \cos \Delta\xi$; the distance of a profile from the mean profile grows with χ while its shape concentration does not; the wage return to χ has no isotropic component; and the directional return is a single first harmonic. One structure generates all four. We posit that the capability content of work is, to first order, a two-dimensional plane in capability space,

$$v_o = \bar{v} + \chi_o \left(\cos \xi_o u_1 + \sin \xi_o u_2 \right) + \varepsilon_o, \quad (1)$$

where v_o is the occupation’s capability profile, $\bar{v}$ the mean profile, and $u_1 \perp u_2$ the two poles of the plane. The task disk is the polar map of this plane: χ is amplitude, ξ is phase. Under the plane, the cosine law for deviation similarity is exact, with noise attenuating it

to $a \cos \Delta\xi$ and more so for central pairs; the deviation magnitude is χ itself; the profile shape is free of χ ; and distance on the disk is, up to the common pole norm, distance in capability space. This is the structural postulate of the model. Every field the model carries over the disk, requirements and prices alike, inherits its functional form from it; higher-order structure is texture, measured and reported in Section 4 but carried outside the theory. Section 4 also tests the postulate directly: the leading two components hold roughly three quarters of the capability deviation variance, with a sharp gap to the third, and the leading score plane aligns with the disk coordinates.

The market enters through one further primitive: capabilities are priced by a linear functional, so the log price of a unit of work is affine in the capability content of its location, $\ln \Pi(\mathbf{r}) = \kappa + \sum_k p_k q_k(\mathbf{r})$. Under the plane this implies that the depth return is a single first harmonic, $\chi (p \cdot u_1 \cos \xi + p \cdot u_2 \sin \xi)$, and that the isotropic depth term is zero; both are confirmed in the estimation (Section 4). We summarize the price of skill as a field $\Pi(\mathbf{r})$ over the disk, the price the market pays for a unit of work performed at task location $\mathbf{r} = (\xi, \chi)$. It is the reduced form of the directional wage equation of Storck and Andersson (2026), a first-harmonic field on the circle whose radial term is modulated by direction,

$$\ln \Pi(\mathbf{r}) = m_0 + m_1 \cos \xi + m_2 \sin \xi + \chi (m_3 + m_4 \cos \xi + m_5 \sin \xi). \quad (2)$$

The direction-modulated radial term is what the plane generates, and it is what makes depth rewarded in the north and unrewarded in the south. The level terms m_1, m_2 are measured structure outside the plane: the part of the wage that does not run through depth-built capabilities, retained in the reduced form and quantified as texture in Section 4. Π is estimated once, on the cross-section of occupations, and treated as fixed structure here. Calling it a reduced form is precise: the coefficients are identified at occupational centroids, where wages are observed, while the model reads the field at task locations, where production happens. The step from the one to the other is testable, because the model prices an occupation as its bundle, eq. (3) below, not as the field at its centroid; Section 4 shows that the bundle-integrated wage reproduces observed wages as well as the centroid evaluation does, and quantifies the gap between the two.

The field is identified on the occupational support, which extends to $\chi = 0.752$, the largest centroid radius in the estimation sample; tasks extend to $\chi = 1$. We treat (2) as structure over the full disk, so evaluating it at peripheral task locations is a linear extrapolation in χ beyond the occupational support. The task layer is defined to the rim, and the field follows it; Section 4 reports the quantitative weight of the extrapolated region.

Productivity is isotropic. Storck and Andersson (2026) does not find that workers are more productive in the north. It finds that the market pays more for the capabilities

concentrated there. Productivity belongs to the worker and is symmetric; the price of that work is directional. The directional structure therefore lives in Π , not in a productivity kernel.

3.3 Tasks, Occupations, and Wages

A task is a unit of work activity that yields a unit of output. Because one task is one unit of output, no productivity magnitude varies across tasks. What varies over the disk is how much labor is present and what that labor is paid.

An occupation o is a bundle: a weight measure $b_o(\mathbf{r})$ over task space, normalized to one, weighted by the O*NET task-relevance ratings, the same weights that define the occupational centroids in Storck and Andersson (2026). Its centroid $\mu_o = \int \mathbf{r} b_o(\mathbf{r}) d\mathbf{r}$ is a summary of the bundle, not the object itself. The occupation's wage is the price of its bundle,

$$w_o = \int \Pi(\mathbf{r}) b_o(\mathbf{r}) d\mathbf{r}, \quad (3)$$

the average price of skill over the tasks the occupation performs. The wage is therefore not a property of the occupation but an aggregate of the price field over its current bundle, and it moves when the bundle is rewritten. Labor density on the disk is

$$n(\mathbf{r}) = \sum_o L_o b_o(\mathbf{r}), \quad (4)$$

where L_o is employment in occupation o , the employment-weighted task coverage of the disk.

3.4 Technology and the Operated Regime

A technology K is a field $\phi_K(\mathbf{r})$, a Gaussian centered at p_K with reach z_K ,

$$\phi_K(\mathbf{r}) = A_K \exp\left(-\frac{1}{2} \left(\frac{\|\mathbf{r} - p_K\|}{z_K}\right)^2\right), \quad (5)$$

the effectiveness with which capital can do the work at $\mathbf{r}$. Capital does not perform a task autonomously. At a location the worker remains and the output is unchanged; what the technology changes is the share of the location's micro-tasks that capital bears. The position p_K fixes where in task space the technology bites, and the reach z_K fixes how far it extends: a narrow z_K is a specialized technology, a broad z_K a general-purpose one. Producing the crossed micro-tasks more cheaply with capital is the productivity effect of automation in the sense of Acemoglu and Restrepo (2018), here made spatial.

A technology also has a character. The same effectiveness can replace the worker or assist her, and these have different consequences. We carry the character as a scalar

$s_K \in [0, 1]$, a fourth attribute of the technology beside p_K , z_K , and A_K . At $s_K = 1$ the technology is purely replacing, at $s_K = 0$ purely augmenting. The effectiveness splits into a replacing part $s_K \phi_K$ and an augmenting part $(1 - s_K) \phi_K$, which act on the same location through different channels. The augmenting part assists rather than replaces; we treat it in Section 3.8.

Whether capital bears a micro-task is a question of price, and only the replacing part competes with the worker on it. Producing the task with labor costs $\Pi(\mathbf{r})$, the price of the work. Producing it with the replacing part of capital costs the rental R against that part's effectiveness, an effective cost $R/(s_K \phi_K(\mathbf{r}))$. Capital is the cheaper producer when

$$s_K \phi_K(\mathbf{r}) > \frac{R}{\Pi(\mathbf{r})}. \quad (6)$$

The comparison is between the rental of capital and the price of labor at the location, weighted by the technology's effectiveness there. The threshold $R/\Pi(\mathbf{r})$ is lowest where $\Pi(\mathbf{r})$ is highest, so capital is profitable first at the most highly priced locations, and a broad, inexpensive technology can sustain it well beyond its core. Capital moves first against the dear work. This is the spatial counterpart of the threshold in Acemoglu and Restrepo (2018): where they automate the lowest task on a single scale, the field takes the highest-priced positions wherever they lie on the disk.

A location $\mathbf{r}$ is not a single task. It is a neighborhood holding a continuum of micro-tasks, and the effectiveness ϕ_K varies across them, so (6) is not true or false for the whole cell at once. It holds for a fraction of the micro-tasks and fails for the rest. We write that fraction as $a(\mathbf{r})$, the degree to which capital bears the work at $\mathbf{r}$,

$$a(\mathbf{r}) = \frac{1}{1 + \exp(-(s_K \phi_K(\mathbf{r}) - R/\Pi(\mathbf{r}))/\tau)}, \quad (7)$$

where τ sets the width of the margin, the dispersion of (6) across the micro-tasks of the cell. As $\tau \rightarrow 0$ the cell collapses to a single task and $a(\mathbf{r})$ becomes a step, zero below the threshold and one above it; in that limit the operated region is a sharp partition of task space and the margin is a boundary in space. For $\tau > 0$ the margin lies inside each cell, and $a(\mathbf{r})$ is the share of its micro-tasks on the operated side of (6).

The replacing part sets the takeover; the augmenting part does not enter it. At $s_K = 1$ the threshold (6) and the operated share (7) are those of pure replacement, and the displacement, wage, and density results that follow hold unchanged. The replacement model is the case $s_K = 1$.

Displacement is then a loss of priced contribution, not a loss of the task. The worker remains at $\mathbf{r}$ and the output is unchanged; what shrinks is the share of that output the worker is paid for. The priced contribution of the cell is $\Pi(\mathbf{r})$ scaled by the retained share $1 - a(\mathbf{r})$. The mass displaced from occupation o is the operated share integrated over its

own bundle,

$$D_o = \int b_o(\mathbf{r}) a(\mathbf{r}) d\mathbf{r}, \quad (8)$$

and its retained wage is $\int \Pi(\mathbf{r}) b_o(\mathbf{r}) (1 - a(\mathbf{r})) d\mathbf{r}$. The aggregate displaced labor mass is $\Delta\Gamma^D = \sum_o L_o D_o$. An occupation can carry substantial D_o without any region having become operated as a whole, because the displacement is spread across its bundle rather than concentrated in a zone it owns. This is the difference between an occupation disappearing and an occupation that remains but is paid for less of what it does.

Because $a(\mathbf{r})$ varies within the bundle, the margin cuts inside the occupation rather than between regions. Some of an occupation's own tasks are borne by capital while others remain with the worker. This is the restructuring of work within an occupation, the unbundling and rebundling of Brynjolfsson, Mitchell, and Rock (2018), expressed as a field. The reach of the technology, p_K and z_K , together with the price field Π , determines which tasks within each bundle cross the margin.

3.5 Reinstatement: The Bundle Operator

Reinstatement returns a fraction γ of the displaced mass as new tasks, and the new tasks must attach to occupations. Free tasks have no role in the model. We therefore need a rule for where new tasks appear and which occupations can take them.

Where they appear. The gradient of the augmentation field,

$$\|\nabla\phi_K(\mathbf{r})\| = \frac{1}{z_K^2} \|\mathbf{r} - p_K\| \cdot \phi_K(\mathbf{r}), \quad (9)$$

peaks at distance z_K from p_K . New tasks are seeded on this boundary ring. With $\hat{g}(\mathbf{r}) = \|\nabla\phi_K(\mathbf{r})\| / \int \|\nabla\phi_K\|$ the normalized ring density and $M = \gamma \Delta\Gamma^D$ the seeded mass, the seeded density is $s(\mathbf{r}) = M \hat{g}(\mathbf{r})$.

Which occupations can take them. A new task at $\mathbf{r}$ requires a profile of capabilities. The requirement fields are not a modeling choice: under the postulate of Section 3.2 they are the plane itself, read per capability cluster,

$$q_k(\xi, \chi) = \bar{v}_k + \chi (u_{1,k} \cos \xi + u_{2,k} \sin \xi), \quad (10)$$

for $k \in \{\text{S1, S2, A1, A2}\}$, with coefficients fitted once to occupation-level cluster intensities (Section 4); a measured variant that adds the level harmonics and an isotropic depth term is carried as a sensitivity. The fields read the requirement at a location from the capabilities of the occupations positioned there, a revealed-requirement assumption. An occupation can attach the task readily when it already holds those capabilities, and only with difficulty when the task requires a capability the occupation lacks. The relevant

quantity is the deficit, the positive shortfall, not a symmetric distance:

$$\delta_o(\mathbf{r}) = \sum_{k \in \mathcal{K}} v_k \max(q_k(\mathbf{r}) - q_{o,k}, 0), \quad (11)$$

where $q_k(\mathbf{r})$ is the level of capability cluster k required at $\mathbf{r}$, $q_{o,k}$ is the occupation's level on cluster k , v_k is the weight of cluster k , and $\mathcal{K}$ is the set of priced clusters. Attaching tasks that use less of the occupation's capabilities is free. Attaching tasks that demand more, or a different kind, is costly.

The weights are not free parameters: they are the components of the price functional of Section 3.2, $v_k = p_k$, measured by the mediation analysis of Storck and Andersson (2026) and replicated on the frozen inputs (Section 4): large for the socio-cognitive skill cluster S1, small for the technical cluster S2, not significant for the ability clusters A1 and A2. The restriction to priced clusters follows from the same identification, not from a significance convention: a capability the functional does not price lies in its kernel, and a shortfall in it is no economic barrier in the sense the operator formalizes; the sign restriction $v_k \geq 0$ excludes carrying noise estimates as theory. Empirically $\mathcal{K} = \{\text{S1}, \text{S2}\}$; the all-cluster variant is a sensitivity. A single functional thus does two jobs. It prices the work a location demands, and it gates what a bundle can absorb. New tasks attach freely only where the deficit is small, that is, where the priced capability content does not exceed what the occupation already holds. Reinstatement can therefore broaden a bundle but cannot lift its priced content. The wage cannot jump. A large jump would require a capability the occupation does not have, and that is not reinstatement within an occupation but the seed of a new one.

In the one-dimensional framework, comparative advantage is a monotone schedule: labor is strongest in the highest-indexed tasks, and new tasks are added there by construction (Assumption 1 of Acemoglu and Restrepo (2018)). The deficit gate is the geometric counterpart. It does not posit a single direction in which labor is strong. It lets a new task attach where the occupation already holds the priced capability the task requires, wherever that lies on the disk.

Attachment and allocation. The occupation's readiness to attach a task at $\mathbf{r}$ is

$$e_o(\mathbf{r}) = \exp(-\delta_o(\mathbf{r})/\ell), \quad (12)$$

between zero and one, where ℓ is a single global scale. The total attachment capacity at $\mathbf{r}$ is $C(\mathbf{r}) = \sum_o L_o e_o(\mathbf{r})$, the labor present weighted by its readiness. A saturating function $\Phi(C) = C/(1 + C)$ gives the bound fraction of the seeded mass. Occupation o receives the reinstated density

$$\iota_o(\mathbf{r}) = s(\mathbf{r}) \Phi(C(\mathbf{r})) \frac{L_o e_o(\mathbf{r})}{C(\mathbf{r})}, \quad (13)$$

its share of the bound mass in proportion to its readiness.

The bundle operator. The post-technology bundle removes displaced weight and adds reinstated weight,

$$b_o^{\text{post}}(\mathbf{r}) = b_o(\mathbf{r}) (1 - a(\mathbf{r})) + \frac{\iota_o(\mathbf{r})}{L_o}. \quad (14)$$

The same field ϕ_K drives both terms. It deactivates through $a(\mathbf{r})$, high near the core, and seeds through $\hat{g}(\mathbf{r})$, high on the ring. The bundle is rewritten continuously as the technology intensifies. This is the geometric form of occupational redefinition (Brynjolfsson, Mitchell, and Rock 2018): the occupation does not vanish, its bundle is redrawn.

The operator does not preserve the bundle's mass. Integrating (14) gives $\int b_o^{\text{post}} = 1 - D_o + B_o/L_o$, where $B_o = \int \iota_o d\mathbf{r}$ is the seeded mass the occupation binds. We do not renormalize. Renormalization would assert that the freed time is spread proportionally over the remaining tasks, a behavioral claim that contradicts the fixed output value of a place (Section 3.7): there is no slack demand for additional human work at the retained locations. It would also double-count the model's actual channel for freed labor, the re-sorting of the worker layer. The mass deficit is that channel's input: $D_o - B_o/L_o$ is, per worker, exactly the labor the occupation releases to the worker layer. After the operator, the bundle is a measure of performed and paid content rather than a unit weight, and the wage (3) is the priced content of that measure, not an average over it; the two readings coincide only pre-technology.

Wages and the Turing trap. The wage change is the price-weighted change in the bundle,

$$\Delta w_o = \int \Pi(\mathbf{r}) [b_o^{\text{post}}(\mathbf{r}) - b_o(\mathbf{r})] d\mathbf{r} = \int \Pi(\mathbf{r}) \left[-b_o(\mathbf{r}) a(\mathbf{r}) + \frac{\iota_o(\mathbf{r})}{L_o} \right] d\mathbf{r}. \quad (15)$$

The first term is negative and falls on high- Π tasks; the second adds tasks that, by the deficit gate, are mostly low- Π . Default reinstatement therefore carries a downward wage bias: it strips priced work and refills with unpriced work. The sign of Δw_o is the contest between where $a(\mathbf{r})$ removes weight and where ι_o adds it, read against the price gradient $\nabla \Pi$. In the directions where depth is priced, the strip falls on high- Π tasks and the refill cannot replace them.

Density. Summing the bundle change over occupations gives the evolved task density,

$$\rho_1(\mathbf{r}) - \rho_0(\mathbf{r}) = - \sum_o L_o b_o(\mathbf{r}) a(\mathbf{r}) + \sum_o \iota_o(\mathbf{r}) + u(\mathbf{r}), \quad (16)$$

with $u(\mathbf{r})$ defined in Section 3.6. The added density integrates to $\gamma \Delta \Gamma^D$, so γ keeps its interpretation as the fraction of displaced mass returned. Setting $\gamma = 0$ recovers pure

automation; $\gamma = 1$ implies full reinstatement.

3.6 Unbound Tasks and Candidate Occupations

Not all seeded mass attaches. Where every occupation has a large deficit, $C(\mathbf{r})$ is small, $\Phi(C(\mathbf{r}))$ is near zero, and the seeded mass stays unbound:

$$u(\mathbf{r}) = s(\mathbf{r}) [1 - \Phi(C(\mathbf{r}))]. \quad (17)$$

u is the density of new tasks that no existing occupation can perform. It is an output of the model, not an agent. The bound mass $\sum_o \iota_o$ and the unbound mass u together integrate to the seeded $\gamma \Delta \Gamma^D$, so the closed-occupation accounting holds.

The map of u answers two distinct questions. Where unbound tasks are created is u itself, the ring regions that are active but that no bundle can absorb. Where new occupations may emerge is where u concentrates, a connected region carrying enough mass to support an occupation, not merely a high value at a point. A diffuse u along the ring is friction. A clump is a candidate.

The price field tells the two apart. Unbound mass in a high- Π region is latent skilled work awaiting a bearer, the kind of region from which qualified new occupations have historically grown. Unbound mass in a low- Π region is work no occupation wants and no worker is paid well for, closer to precarious friction than to a new occupation. We color candidate regions by Π and read opportunity against spillage. We do not instantiate new occupations: the model shows where one may appear, not that one does.

The same deficit that gates attachment draws the boundary between two literatures. New tasks are absorbed by existing occupations where the priced deficit is small, which is redefinition in the sense of Brynjolfsson, Mitchell, and Rock (2018). They remain unbound where the deficit is large for every occupation, which is the new work of Acemoglu and Restrepo (2019) that no existing occupation owns. One operator decides which, through the deficit.

3.7 Employment and Mobility

The worker layer sets employment L_o and closes the model. We solve it as a static spatial-sorting equilibrium and read the effect of a technology as the comparative static between two such equilibria, one with the original bundles and one with the bundles rewritten by (14). A law of motion in time can be added on top of this rest point; we do not need it for a statement about the new equilibrium.

Two factors share each place. A location is produced by human and capital input as perfect substitutes, in the manner of Acemoglu and Restrepo (2018). The bundle operator

splits each occupation's coverage of a place into a human part $h_o(\mathbf{r}) = b_o^{\text{post}}(\mathbf{r})$, the work the worker performs, and a capital part $k_o(\mathbf{r}) = b_o(\mathbf{r}) a(\mathbf{r})$, the work capital bears. Total task-work at $\mathbf{r}$ is

$$n(\mathbf{r}) = \sum_o L_o (h_o(\mathbf{r}) + k_o(\mathbf{r})) = \sum_o L_o b_o(\mathbf{r}) + \sum_o \iota_o(\mathbf{r}), \quad (18)$$

the pre-technology density (4) plus the bound seeded mass; unbound mass performs no work and does not enter. The operated share enters once, inside h_o ; no further $(1 - a)$ factor is applied downstream. Pre-technology, $h_o = b_o$, $k_o = 0$, and n reduces to (4). Output value is

$$V(\mathbf{r}) = \Pi(\mathbf{r}) n(\mathbf{r})^\beta, \quad \beta \in (0, 1), \quad (19)$$

concave in task-work through the congestion exponent β , and unchanged by who performs the work. We hold the output value of a place fixed under the human-capital split. This keeps the framework partial: it carries no channel for goods prices or growth, which would require modeling the wider economy, and we leave that to future work. Congestion acts on n , the total task-work present, regardless of which factor bears it; the takeover redistributes the work between h_o and k_o without changing their sum at a location. Because the pre-technology bundle b_o is normalized, a broad task radius spreads it thin and contributes little to $n(\mathbf{r})$, so a concentrated occupation presses harder on the places it occupies than a diffuse one does.

Returns to the two factors. Human and capital are paid the marginal product of task-work, $\beta \Pi(\mathbf{r}) n(\mathbf{r})^{\beta-1}$, each on the part of the work it performs: the human side on $\sum_o L_o h_o(\mathbf{r})$, the capital side on $\sum_o L_o k_o(\mathbf{r})$. Their payments sum to $\beta V(\mathbf{r})$; the residual $(1 - \beta)V(\mathbf{r})$ is the congestion rent to the fixed local factor.

Occupation value and equilibrium. A worker carries the whole bundle and earns the marginal product integrated over the content she performs,

$$W_o = \beta \int h_o(\mathbf{r}) \Pi(\mathbf{r}) n(\mathbf{r})^{\beta-1} d\mathbf{r}, \quad (20)$$

the congestion- and takeover-adjusted counterpart of the gross bundle price w_o of (3): the takeover enters through the stripped part of h_o , congestion through $n^{\beta-1}$, the price through Π . Occupation values are coupled, because a location's density sums over every occupation that covers it, so occupations that overlap in task space press on one another's returns. Labor re-sorts from the pre-technology distribution L_o^0 , drawn toward value and damped by the bundle distance $d(o, o') = \|\mu_o - \mu_{o'}\|$, the Euclidean distance between occupational centroids, which Storck and Andersson (2026) ties to the transferability of human capital in the sense of Gathmann and Schönberg (2010). Under the capability plane the metric is

induced rather than chosen: with orthogonal poles of equal norm, distance on the disk is proportional to distance in capability space, so the damping measures the capability gap a mover must close. The allocation is a logit over value net of the mobility cost from origin,

$$L_o = L^{\text{tot}} \frac{L_o^0 \exp((W_o - c d(o, o^0))/\kappa)}{\sum_{o'} L_{o'}^0 \exp((W_{o'} - c d(o', o^0))/\kappa)}, \quad (21)$$

a fixed point because W_o depends on $n(\mathbf{r})$, which depends on $\{L_o\}$. The distance term bites in a static solution because the re-sorting is measured from an origin, not from nowhere. Total labor is fixed at L^{tot} ; letting it respond to aggregate value is a later extension.

The labor share. With $H(\mathbf{r}) = \sum_o L_o h_o(\mathbf{r})$ the human task-work, total labor payment is $\int \beta \Pi(\mathbf{r}) H(\mathbf{r}) n(\mathbf{r})^{\beta-1} d\mathbf{r}$ and total factor payment to human and capital is $\int \beta \Pi(\mathbf{r}) n(\mathbf{r})^\beta d\mathbf{r}$, so the labor share of factor income is

$$\Lambda = \frac{\int \Pi(\mathbf{r}) H(\mathbf{r}) n(\mathbf{r})^{\beta-1} d\mathbf{r}}{\int \Pi(\mathbf{r}) n(\mathbf{r})^\beta d\mathbf{r}}. \quad (22)$$

The congestion exponent β cancels: it governs the spatial sort but not the share. The share is the price- and density-weighted mean of the human share $H(\mathbf{r})/n(\mathbf{r})$, which is one pre-technology; automation lowers it by raising a where (6) is crossed, and with no reinstatement it reduces to the mean of $1 - a(\mathbf{r})$. Solving (21) once with the original bundles and once with the rewritten bundles gives two equilibria, and the difference in Λ is the effect of the technology. The displacement of labor onto fewer tasks then appears not at the automated places, where a is high and density falls, but at their destinations, where the re-sorting raises n and congestion bites. The bunching of Acemoglu and Restrepo (2018) is recovered as an equilibrium effect of the sort rather than assumed locally, and whether the share falls hardest in priced or unpriced directions is decided by where a rises against the price field.

3.8 Augmentation and the Character of the Technology

The augmenting part $(1 - s_K) \phi_K$ does not take the task. It lowers the labor the worker needs to produce it. The human labor requirement at $\mathbf{r}$ is scaled by

$$\frac{1}{1 + (1 - s_K) \phi_K(\mathbf{r})}, \quad (23)$$

so the worker covers the same output with less labor and the output of the location is unchanged. The labor this releases does not leave the model. It enters the mobile pool and re-sorts through the allocation (21), alongside the labor freed by replacement. Where it lands is set by occupation value and mobility cost, not by the location it left.

The two parts differ in three ways. *Completeness.* Replacement takes the whole priced contribution of the crossed micro-tasks, more of the bundle the higher s_K is and sooner the lower R is, because both widen the region where (6) holds. Augmentation never takes the task; it scales the requirement by a bounded factor and leaves the work with the worker. *Feedback.* Replacement is self-damping in its own direction: as priced mass passes to capital the location’s average price falls, the threshold R/Π rises, and further takeover slows, the brake of Acemoglu and Restrepo (2018) inherited through the price field. Augmentation has no such brake, because it moves no priced mass to capital and does not change Π . *Destination.* Replacement seeds reinstatement on the gradient ring, so part of the displaced mass returns near where it was lost, gated by the deficit. Augmentation seeds nothing: reinstatement is a fraction of displaced mass, and the augmenting end displaces nothing, so $M = 0$ and $\iota_o = 0$ there. The freed labor only re-sorts, and its outcome is set entirely by the priced locations it can reach. The augmenting end is a statement about destinations.

The main model takes s_K near one, where capital replaces and reinstatement follows. The pure augmenting end $s_K = 0$ is the contrast: no takeover, no seeding, only labor freed and re-sorted.

Two bounds on this treatment. The character s_K is a scalar, uniform over the reach of the technology. A technology that replaces in one direction and assists in another is a direction-dependent character, $s_K(\mathbf{r})$ as a field, which we leave as an extension. The augmenting end is also where the model most strains its static closure: separating the labor present at a location from the labor it requires decouples congestion from output, which the worker layer of Section 3.7 does not yet carry. We hold the main results at s_K near one, where labor present and task-work coincide and the accounting below is exact, and develop the destination problem of freed labor in later work on occupational mobility.

4 Empirical Implementation and Calibration

This section measures the frozen structure of the model: the capability plane, the price field, the bundle pricing it implies, the components of the price functional, and the wage wedge. Every number reported here is the output of a script in the replication repository, and the inputs are frozen copies of the Paper 1 exports with recorded provenance. The technology parameters, the only objects a scenario is allowed to vary, are calibrated in Section 4.7.

4.1 Data and frozen inputs

All inputs are vendored from a pinned commit of the Storck and Andersson (2026) repository by a single freezing script, which is the model’s only contact surface with the

upstream project; every analysis script reads exclusively from the frozen copies, and a manifest records the source commit, the encoder run, and per-file hashes and derivation recipes. The frozen inputs are the occupation and task polar coordinates of the reference encoder run (878 occupations; task coordinates and relevance ratings defining the bundles b_o), BLS OEWS May 2023 median hourly wages, the frequency-weighted mean required level of education per occupation, the capability cluster intensities $q_{o,k}$, and the occupation-by-descriptor matrices of raw levels for the 35 O*NET Skills and 52 Abilities. Merging coordinates with positive median wages and education reproduces the Paper 1 estimation sample of $N = 785$ occupations; the capability-side estimations, which do not require wages, use all 878.

4.2 The capability plane

The postulate of Section 3.2 is testable directly, and the test uses data that enter no step in the construction of the coordinates: the disk is built from task text alone, while the descriptor matrices record surveyed capability levels. For each descriptor family we compute the principal components of the deviations $v_o - \bar{v}$ and ask two questions: how much of the deviation variance lies in a two-dimensional structure, and whether that structure is the disk.

The answers are, respectively, three quarters and yes to first order. For Skills, the leading two components hold 73.3 percent of the deviation variance (46.2 and 27.0), with a sharp gap to the third (5.7); for Abilities, 75.4 percent (62.7 and 12.7, against 5.8 for the third). The leading score plane aligns with the disk coordinates $(x, y) = (\chi \cos \xi, \chi \sin \xi)$ with canonical correlations 0.93 and 0.72 for Skills and 0.89 and 0.56 for Abilities; the disk coordinates are recovered from the scores with $R^2 = 0.85$ (x) and 0.55 (y) for Skills, 0.64 and 0.49 for Abilities. The capability content of work is two-dimensional to first order, and the task disk is an imperfect but unmistakable polar map of that plane: the first canonical direction is nearly identified, the second is recovered more coarsely.

The plane read per capability cluster gives the requirement fields of eq. (10), fitted by OLS on the 878 occupations (Table 1). The pole structure matches the compass of Storck and Andersson (2026): the socio-cognitive clusters S1 and A1 point northeast (52° and 54°), the technical cluster S2 northwest (151°), the physical cluster A2 west (189°). The exact plane form, three parameters per cluster, fits nearly as well as the measured six-parameter variant: adding the level harmonics and the isotropic depth term raises R^2 by at most 0.009 (Table 1, last column). The theory costs almost nothing relative to the unrestricted first-harmonic fit.

What the plane does not generate is recorded in a texture ledger (Table 2), with the same discipline as the wage wedge: measured once, carried outside the theory, switched on in sensitivity runs. The entries are the level harmonics of the price field (the eastward

Table 1: The capability plane read per cluster, eq. (10): OLS on $N = 878$ occupations, HC3 standard errors in parentheses. Pole direction is $\text{atan2}(u_{2,k}, u_{1,k})$, amplitude $\|(u_{1,k}, u_{2,k})\|$. ΔR^2 is the gain from the measured variant (adding level harmonics and an isotropic depth term).

Cluster	$\bar{v}_k$	$u_{1,k}$	$u_{2,k}$	Pole	R^2	ΔR^2 (meas.)
S1	2.682 (0.011)	+1.094 (0.027)	+1.421 (0.048)	52°	0.737	+0.001
S2	1.366 (0.018)	−2.219 (0.048)	+1.228 (0.080)	151°	0.686	+0.008
A1	3.198 (0.009)	+0.952 (0.024)	+1.309 (0.041)	54°	0.748	+0.000
A2	1.573 (0.016)	−1.539 (0.038)	−0.244 (0.065)	189°	0.571	+0.009

gradient $m_1 = +0.19$), the east–west level second harmonic of wages (± 9 percent at fixed χ), the cluster-level gains of Table 1, and the second-harmonic remainders of the requirement fields, which are statistically detectable at $N = 878$ but add at most 0.030 (0.011) of R^2 in level (interaction).

Table 2: Texture ledger: measured structure outside the capability plane. Each entry is estimated once and held fixed; regime runs are reported with and without the texture-bearing variants.

Texture	Magnitude	Where measured
Wage level gradient (east)	$m_1 = +0.192$, dir. 25°	price-field estimation
Wage level 2nd harmonic	$\pm 9\%$ E–W at fixed χ	price-field estimation
q_k level harmonics + isotropic depth	$\Delta R^2 \leq 0.009$	plane estimation
q_k 2nd harmonics (level; interaction)	$\Delta R^2 \leq 0.030$; ≤ 0.011	plane estimation

4.3 The price field

The coefficients of eq. (2) are estimated by OLS of log median hourly wage on the polar regressors over the $N = 785$ sample, with HC3 standard errors. The estimation first replicates the geometry-only regression of Storck and Andersson (2026) exactly ($\cos \xi + 0.212$, $\sin \xi + 0.389$, $\chi - 0.046$, $R^2 = 0.492$), pinning the sample, and then adds the interaction terms. The estimated field is

$$\ln \Pi = \underset{(0.036)}{3.420} + \underset{(0.048)}{0.192} \cos \xi + \underset{(0.051)}{0.089} \sin \xi + \chi \left(\underset{(0.090)}{-0.065} + \underset{(0.112)}{0.008} \cos \xi + \underset{(0.134)}{0.891} \sin \xi \right),$$

$R^2 = 0.523$. The two predictions of the plane-with-linear-pricing structure hold: the isotropic depth term is indistinguishable from zero ($m_3 = -0.065$, $p = 0.47$), and the depth return is a single first harmonic, peaking at 89.5° with amplitude 0.891 and positive for $\xi \in (3.6^\circ, 175.3^\circ)$. The balanced second-harmonic test leaves the interaction terms jointly insignificant conditional on the level terms ($p = 0.44$); the naive test that omits the level terms rejects spuriously, because $\chi \cos 2\xi$ is nearly collinear with $\cos 2\xi$ (correlation

0.93). The second-order structure that exists is a level effect of about ± 9 percent along the east–west axis ($\cos 2\xi$: $+0.089$, $p < 0.001$), entered in the texture ledger.¹ An employment-weighted estimation leaves the depth harmonic intact ($m_5 = +0.848$, $R^2 = 0.704$).

4.4 From centroids to bundles

The reduced-form step of Section 3.2 is from coefficients identified at occupational centroids to a field read at task locations, and eq. (3) makes it testable: the model prices an occupation as its bundle, not as the field at its centroid. The bundle-integrated wage predicts observed log wages with $R^2 = 0.510$ against 0.523 for the in-sample centroid evaluation (prediction R^2 without refitting; correlation 0.717), and the Jensen gap $\ln w_{\text{bundle}} - \ln \Pi(\mu_o)$ is small (mean $+0.018$, sd 0.049, range -0.168 to $+0.195$). Bundle dispersion matters little for pricing in the cross-section, which is what licenses estimating the field at centroids and operating it at tasks.

The same step crosses the support boundary stated in Section 3.2: the field is identified on centroids with $\chi \leq 0.752$, while tasks extend to the rim. The extrapolated region carries on average 7.3 percent of an occupation’s bundle mass (median zero, maximum 66.7 percent; 7.5 percent of tasks unweighted). Its quantitative weight in pricing is negligible: clipping the field at the support boundary changes bundle log wages by $+0.0007$ on average (sd 0.0044, maximum 0.038) and moves the prediction R^2 from 0.5097 to 0.5101, while the extrapolated price at the northern rim exceeds the boundary value by 23 percent. The model uses the extrapolated field, consistent with the task layer being defined to the rim; the clipped variant is a sensitivity.

4.5 The price functional and the gate weights

The components p_k of the price functional, which the deficit gate of eq. (11) uses as weights v_k , are measured by replicating the global mediation regression of Storck and Andersson (2026) on the frozen inputs: log wage on the direction terms, χ , mean required education, and the four cluster intensities ($N = 785$, HC3, $R^2 = 0.683$). The replication is exact against the published values: $\chi -0.079$ ($p = 0.24$; published -0.08 , $p = 0.24$), S1 $+0.328$ ($p < 0.001$; published $+0.33$), S2 $+0.049$ ($p = 0.020$; published $+0.05$, $p = 0.02$), A1 -0.030 and A2 $+0.010$ (both insignificant). The direction terms are part of the replicated specification; the variant without them does not replicate. The gate accordingly runs on the priced clusters, $\mathcal{K} = \{\text{S1}, \text{S2}\}$ with $v = (0.328, 0.049)$, as derived in Section 3.5; the all-cluster variant is a sensitivity in the regime runs.

¹Since m_4 is indistinguishable from zero, the depth-priced core of the field is, in Cartesian coordinates, $\ln \Pi \approx m_0 + 0.89 y$ plus the level terms: a single linear coordinate of the plane. We note this in passing; the model carries the full reduced form.

4.6 The wage wedge

The price field accounts for roughly half of the wage cross-section; the residual variance share is 0.477, with weak but significant spatial autocorrelation (Moran's I +0.080, permutation $p = 0.001$). The residual is not capability structure in disguise: the four cluster intensities explain 9.7 percent of it, while Job Family fixed effects explain 27.9 percent (33.4 jointly with the intensities, for a total R^2 of 0.682 against log wages). What remains above the field is mostly classification-level wage setting, which is precisely what the theory layer excludes and the measurement layer records. The occupation wedge is $\eta_o = \ln w_{\text{obs},o} - \ln w_{\text{bundle},o}$, and the family wedge η_g is its unweighted family mean, ranging from -0.261 (Personal Care and Service) to $+0.372$ (Management) and absorbing 30.2 percent of the occupation-level wedge variance (employment-weighted means correlate at 0.948). The wedge is measured once on this cross-section and held fixed under technology shocks. It enters the operated regime, when switched on, through the effective price $\Pi_{\text{eff}} = e^{\eta_g} \Pi$, so that dearer families cross the takeover margin sooner; the regime runs of Section 5 are reported with and without it.

4.7 Technology calibration

[TODO: Kalibrering av teknologiparametrarna (p_K, z_K, A_K, s_K) mot task-nivå-AI-exponering, och av attachment-skalan ℓ mot observerade centroidförskjutningar. Datakälla avgörs i steg D: kandidater som redan ligger på O*NET-tasknivå är Eloundou et al. (2024) och SML (Brynjolfsson, Mitchell, and Rock 2018); Gmyrek et al. (2025) kräver överföring från polska task-statements och prövas endast om de direkta källorna faller. Enheter: R och Π i USD per timme, A_K dimensionslös effektivitet, τ marginbredd i samma enhet som $s_K \phi_K - R/\Pi$.]

5 Illustrative Cases

[TODO: Kondenserad version av Case 1 och Case 2 från Paper 1, med kalibrerade parametrar. Fokus på den spatiala tröskelmekanismen och känslighetsanalysen över z_K . Lägg till en figur över $u(r)$, kartan över obundna uppgifter, färglagd med Π .]

6 Empirical Analysis

[TODO: DiD-specifikation, event study, robusthetscheckar. Baserat på notebook 9-resultaten. Lägg eventuellt till testet av buntomritning: projicera $\Delta\mu_o$ på $\nabla\Pi$ och pröva mot observerade löneförändringar 2019–2024, samma 741 yrken.]

7 Discussion

The central mechanism is the directional pricing of work. Reinstatement is good for employment wherever it lands, because it returns task mass to labor. It is good for wages only where the price field rewards the work it returns. New tasks seeded into the analytical directions, where depth is priced, can raise both employment and the wage. The same new tasks seeded into the service and manual directions raise employment but not the wage, because depth is unpriced there. A reinstatement elasticity that looks protective in aggregate can therefore leave manual workers with more work and no more pay. This is a distributional statement that a one-dimensional model cannot make, because it has no direction in which to place the new work.

Default automation and default reinstatement pull the wage in opposite directions, and the price field decides the contest. The first term of (15) strips the dear tasks, since the operated threshold (6) is crossed first where Π is high. The second term refills with tasks that the deficit gate admits, which are mostly cheap. Left to run, the process bleeds priced work and returns unpriced work, the bundle form of the Turing trap. The same price channel supplies its own brake. As a bundle loses priced mass to capital, its average price falls, which raises the threshold R/Π for the work that remains and slows further takeover in that direction. This is the self-correction of Acemoglu and Restrepo (2018), where a wave of automation lowers the effective cost of the labor that is left and discourages further automation, here inherited through the price field rather than assumed.

The displacement falls within occupations, not between regions, because it lives on the interior margin $a(\mathbf{r})$ rather than on a boundary in space. A bundle spans a range of locations, and the operated share varies across them, so some of an occupation's own tasks cross the margin while others do not. This is why the occupation is the wrong unit and the task is the right one. The point has an empirical counterpart: the automatability of work varies more within occupations than between them (Brynjolfsson, Mitchell, and Rock 2018), which is the same statement as a bundle spanning a range of ϕ_K . The within occupation margin is therefore not a modeling convenience but the level at which the displacement actually occurs.

Exposure and incidence are distinct, and the model keeps them apart. The field ϕ_K is technical exposure, what capital can do at a location. The operated share $a(\mathbf{r})$ is economic incidence, what capital is paid to do there, set by exposure against price through (6). The two need not coincide. A weak correlation between exposure and wages does not weaken the trap, because the trap is a claim about where operation is profitable given exposure, not about where exposure happens to fall. Directed implementation concentrates on the work that is both exposed and dear (Brynjolfsson, Mitchell, and Rock 2018), which is exactly the high- Π high- ϕ_K region where (6) is crossed first. Reading exposure alone would miss this; the price field is what turns exposure into incidence.

The deficit gate is also the boundary between two outcomes that the debate often conflates. Where attachment capacity $C(\mathbf{r})$ is positive, the bundle is redrawn and the occupation is redefined. Where $C(\mathbf{r})$ collapses, the seeded mass stays unbound, and where it clusters it marks a candidate occupation that no existing bundle can hold. Redefinition and the birth of a new occupation are then the two sides of the same gate, separated by whether any occupation already holds the priced capability the new work requires.

Several limitations bound these claims. The worker layer is closed as a static equilibrium, so the results are comparative statics between two rest points rather than a path in time: the adjustment dynamics and the response of total labor supply to aggregate value are left for later work. The output value of a place is held fixed under the human-capital split, which keeps the framework partial and carries no channel for goods prices or growth. New tasks enter as added weight on existing locations rather than as genuinely new points on the disk, which understates the possibility that automation opens directions the geometry does not yet contain. The cluster weights v_k and the acquisition scale ℓ are identified from the cross-section and held fixed, and the wage test of the displacement channel, projecting the centroid shift $\Delta\mu_o$ on $\nabla\Pi$ against observed wage changes, is the sharpest available check on whether the takeover framing is right.

[TODO: Relatera till Autor & Thompson (2025) när referensen verifierats; placera i stycket om riktningsberoende prissättning eller gränzonen.]

References

- Acemoglu, Daron and Pascual Restrepo (2018). “The Race between Man and Machine: Implications of Technology for Growth, Factor Shares, and Employment”. In: *American Economic Review* 108.6, pp. 1488–1542. DOI: 10.1257/aer.20160696.
- (2019). “Automation and New Tasks: How Technology Displaces and Reinstates Labor”. In: *Journal of Economic Perspectives* 33.2, pp. 3–30.
- (2020). “Robots and jobs: Evidence from US labor markets”. In: *Journal of Political Economy* 128.6, pp. 2188–2244.
- Autor, David H, Frank Levy, and Richard J Murnane (2003). “The skill content of recent technological change: An empirical exploration”. In: *The Quarterly Journal of Economics* 118.4, pp. 1279–1333.
- Brynjolfsson, Erik, Tom Mitchell, and Daniel Rock (2018). “What Can Machines Learn and What Does It Mean for Occupations and the Economy?” In: *AEA Papers and Proceedings* 108, pp. 43–47. DOI: 10.1257/pandp.20181019.
- De Vries, Gaaitzen J et al. (2020). “The rise of robots and the fall of routine jobs”. In: *Labour Economics* 66, p. 101885.
- Eloundou, Tyna et al. (2024). “GPTs are GPTs: Labor market impact potential of LLMs”. In: *Science* 384.6702, pp. 1306–1308.

- Gathmann, Christina and Uta Schönberg (2010). “How general is human capital? A task-based approach”. In: *Journal of Labor Economics* 28.1, pp. 1–49.
- Gmyrek, Paweł et al. (2025). *Generative AI and jobs: A refined global index of occupational exposure*. 140. ILO Working Paper.
- Graetz, Georg and Guy Michaels (2018). “Robots at work”. In: *Review of Economics and Statistics* 100.5, pp. 753–768.
- Hampole, Menaka et al. (2025). *Artificial intelligence and the labor market*. Tech. rep. National Bureau of Economic Research.
- Storck, Joakim and Jonatan Andersson (2026). “A polar task-based geometry of work”. In: *Working Paper*.

[TODO: Bibposter att lägga till och verifiera i references.bib: `acemoglu2018race` (Acemoglu & Restrepo, “The Race between Man and Machine: Implications of Technology for Growth, Factor Shares, and Employment”, *American Economic Review* 108(6):1488–1542, 2018, DOI 10.1257/aer.20160696 — verifierad mot fulltext); `brynjolfsson2018mlworkforce` (Brynjolfsson, Mitchell & Rock, “What Can Machines Learn and What Does It Mean for Occupations and the Economy?”, *AEA Papers and Proceedings* 108:43–47, 2018, DOI 10.1257/pandp.20181019 — verifierad mot fulltext); `brynjolfsson2022turing` (Brynjolfsson, “The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence”, *Daedalus* 2022 — ej verifierad mot fulltext, kontrollera volym, nummer och sidor). Kontrollera även att `storck2026geometry`, `gmyrek2025` och `acemoglu2019automation` redan finns. Nyckeln `storck2026mobility` används inte längre (beslut 2026-06-11: Paper 2 citeras inte; buntavståndet $d(o, o')$ definieras självständigt via centroidavstånd med `storck2026geometry`). Lägg till och verifiera `gathmann2010` (Gathmann & Schönberg, “How General Is Human Capital? A Task-Based Approach”, *Journal of Labor Economics* 28(1):1–49, 2010 — samma post som i Paper 1).]